



Chapter 5: Advanced SQL

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Outline

- Accessing SQL From a Programming Language
- Functions
- Triggers
- Recursive Queries
- Advanced Aggregation Features



Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language for at least two reasons

- Not all queries can be expressed in SQL, since SQL does not provide the full expressive power of a general-purpose language.
- Non-declarative actions -- such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface -- cannot be done from within SQL.



Accessing SQL from a Programming Language (Cont.)

There are two approaches to accessing SQL from a general-purpose programming language

- A general-purpose program -- can connect to and communicate with a database server using a collection of functions
 - JDBC (Java)
 - ODBC (C, C++, Go, Python)



JDBC



JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the statement object to send queries and fetch results
 - Exception mechanism to handle errors



JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try (Connection conn = DriverManager.getConnection(
        "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement();
    )
    {
        ... Do Actual Work ....
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```

NOTE: Above syntax works with Java 7, and JDBC 4 onwards.
Resources opened in “try (....)” syntax (“try with resources”) are automatically closed at the end of the try block



JDBC Code (Cont.)

- Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values('77987', 'Kim', 'Physics', 98000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

- Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
    from instructor  
    group by dept_name");  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
        rset.getFloat(2));  
}
```




JDBC SUBSECTIONS

- Connecting to the Database
- Shipping SQL Statements to the Database System
- Exceptions and Resource Management
- Retrieving the Result of a Query
- Prepared Statements
- Callable Statements
- Metadata Features
- Other Features
- Database Access from Python



JDBC Code Details

- Getting result fields:
 - **`rs.getString("dept_name")` and `rs.getString(1)` equivalent if `dept_name` is the first argument of select result.**
- Dealing with Null values
 - `int a = rs.getInt("a");`**
 - `if (rs.isNull()) Systems.out.println("Got null value");`**



Prepared Statement

- ```
PreparedStatement pStmt = conn.prepareStatement(
 "insert into instructor values(?,?,?,?)");

pStmt.setString(1, "88877");
pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance");
pStmt.setInt(4, 125000);
pStmt.executeUpdate();
pStmt.setString(1, "88878");
pStmt.executeUpdate();
```
- **WARNING:** always use prepared statements when taking an input from the user and adding it to a query
  - NEVER create a query by concatenating strings
  - "insert into instructor values(' " + ID + " ', ' " + name + " ', " + "' + dept name + " ', "' balance + ')"
  - What if name is "D'Souza"?



# SQL Injection

- Suppose query is constructed using
  - "select \* from instructor where name = " + name + ""
- Suppose the user, instead of entering a name, enters:
  - X' or 'Y' = 'Y
- then the resulting statement becomes:
  - "select \* from instructor where name = " + "X' or 'Y' = 'Y" + ""
  - which is:
    - ▶ select \* from instructor where name = 'X' or 'Y' = 'Y'
  - User could have even used
    - ▶ X'; update instructor set salary = salary + 10000; --
- Prepared statement internally uses:  
"select \* from instructor where name = 'X\'' or \'Y\' = \'Y\'"
  - **Always use prepared statements, with user inputs as parameters**



# Metadata Features

- ResultSet metadata
- E.g. after executing query to get a ResultSet rs:
  - ```
ResultSetMetaData rsmd = rs.getMetaData();  
for(int i = 1; i <= rsmd.getColumnCount(); i++) {  
    System.out.println(rsmd.getColumnName(i));  
    System.out.println(rsmd.getColumnTypeName(i));  
}
```
- How is this useful?



Metadata (Cont)

- Database metadata

- `DatabaseMetaData dbmd = conn.getMetaData();`

`// Arguments to getColumnns: Catalog, Schema-pattern, Table-pattern,
// and Column-Pattern`

`// Returns: One row for each column; row has a number of attributes
// such as COLUMN_NAME, TYPE_NAME`

`// The value null indicates all Catalogs/Schemas.`

`// The value "" indicates current catalog/schema`

`// The value "%" has the same meaning as SQL like clause`

```
ResultSet rs = dbmd.getColumnns(null, "univdb", "department", "%");
```

```
while( rs.next()) {
```

```
    System.out.println(rs.getString("COLUMN_NAME"),
```

```
                        rs.getString("TYPE_NAME");
```

```
}
```

- And where is this useful?



Metadata (Cont)

- Database metadata

- `DatabaseMetaData dbmd = conn.getMetaData();`

`// Arguments to getTables: Catalog, Schema-pattern, Table-pattern,
// and Table-Type`

`// Returns: One row for each table; row has a number of attributes
// such as TABLE_NAME, TABLE_CAT, TABLE_TYPE, ..`

`// The value null indicates all Catalogs/Schemas.`

`// The value "" indicates current catalog/schema`

`// The value "%" has the same meaning as SQL like clause`

`// The last attribute is an array of types of tables to return.`

`// TABLE means only regular tables`

`ResultSet rs = dbmd.getTables ("", "", "%", new String[] {"TABLES"});`

`while(rs.next()) {`

`System.out.println(rs.getString("TABLE_NAME"));`

`}`

- And where is this useful?



Finding Primary Keys

- DatabaseMetaData dmd = connection.getMetaData();

```
// Arguments below are: Catalog, Schema, and Table
// The value "" for Catalog/Schema indicates current catalog/schema
// The value null indicates all catalogs/schemas
ResultSet rs = dmd.getPrimaryKeys("", "", tableName);
```

```
while(rs.next()){
    // KEY_SEQ indicates the position of the attribute in
    // the primary key, which is required if a primary key has multiple
    // attributes
    System.out.println(rs.getString("KEY_SEQ"),
                       rs.getString("COLUMN_NAME"));
}
```




Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
 - bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
 - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
 - `conn.commit();` or
 - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit.



Other JDBC Features

■ Calling functions and procedures

- `CallableStatement cStmt1 = conn.prepareCall("{? = call some function(?)})");`
- `CallableStatement cStmt2 = conn.prepareCall("{call some procedure(?,?)})");`

■ Handling large object types

- `getBlob()` and `getClob()` that are similar to the `getString()` method, but return objects of type `Blob` and `Clob`, respectively
- get data from these objects by `getBytes()`
- associate an open stream with Java `Blob` or `Clob` object to update large objects
 - ▶ `blob.setBlob(int parameterIndex, InputStream inputStream).`



JDBC Resources

- JDBC Basics Tutorial
 - <https://docs.oracle.com/javase/tutorial/jdbc/index.html>



ODBC



ODBC

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - ▶ open a connection with a database,
 - ▶ send queries and updates,
 - ▶ get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC
- **Please follow the lecture 8 on CSE-302 lab.**



Functions and Procedures



Functions

- Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements.
- These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, or C++.
- The syntax we present here is defined by the SQL standard.
 - Most databases implement nonstandard versions of this syntax.



Declaring SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
CREATE OR REPLACE FUNCTION get_instructor_for_dept(d_name IN varchar2)
RETURN NUMBER
IS d_count NUMBER(11,2);
BEGIN
    SELECT count(*)
    INTO d_count
    FROM instructor
    WHERE dept_name = d_name;
    RETURN d_count;
END;
```

```
SELECT get_instructor_for_dept('Finance') FROM dual;
```

- The function *dept_count* can be used to find the department names and budget of all departments with more that 12 instructors.

```
select dept_name, budget
from department
where get_instructor_for_dept (dept_name ) > 1
```




PL/SQL

DECLARE

```
    dept_name varchar2(20) := 'Finance';  
    item varchar2(20);
```

BEGIN

```
    FOR item IN (SELECT DISTINCT(dept_name) FROM instructor) LOOP  
        IF (item.dept_name = dept_name) THEN  
            dbms_output.put_line('Found ' || item.dept_name || ' department!');  
        END IF;  
    END LOOP;  
END;
```



Triggers



Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - Syntax illustrated here may not work exactly on your database system; check the system manuals



Trigger (Example)

```
• -- create audit table
-- DROP TABLE instructor_insert_audit_log;
CREATE TABLE instructor_insert_audit_log(
    instructor_id varchar(5),
    user_name varchar2(64),
    created_at date
);

-- create trigger
CREATE OR REPLACE TRIGGER instructor_audit
BEFORE INSERT
    ON instructor
    FOR EACH ROW

DECLARE
    username varchar2(10);
BEGIN
    INSERT INTO instructor_insert_audit_log VALUES (:NEW.ID, user, sysdate);
END;

-- test with an insert
INSERT INTO instructor VALUES ('9876', 'me', 'Music', 50000);
SELECT * FROM instructor_insert_audit_log;
-- cleanup
-- DELETE FROM instructor WHERE name = 'me';
```



Trigger to Maintain credits_earned value

```
CREATE OR REPLACE TRIGGER credits_earned
AFTER UPDATE
  OF grade
  ON takes
  FOR EACH ROW
DECLARE
  course_credit NUMERIC(2, 0);
BEGIN
  IF (:NEW.grade <> 'F' AND :NEW.grade IS NOT NULL) AND (:OLD.grade = 'F' OR :OLD.grade IS NULL)) THEN
    DBMS_OUTPUT.PUT_LINE('Updating total credit in student table...');
    SELECT credits INTO course_credit FROM course WHERE course_id = :NEW.course_id;
    UPDATE student SET tot_cred = (tot_cred + course_credit) WHERE id = :NEW.id;
  END IF;
END;
```



Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called *transition tables*) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows

```
CREATE TRIGGER FLIGHTS_DELETE
AFTER DELETE ON FLIGHTS
REFERENCING OLD_TABLE AS DELETED_FLIGHTS
FOR EACH STATEMENT
DELETE FROM FLIGHT_AVAILABILITY WHERE FLIGHT_ID IN
(SELECT FLIGHT_ID FROM DELETED_FLIGHTS);
```



When Not To Use Triggers

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
 - Define methods to update fields
 - Carry out actions as part of the update methods instead of through a trigger



When Not To Use Triggers (Cont.)

- Risk of unintended execution of triggers, for example, when
 - Loading data from a backup copy
 - Replicating updates at a remote site
 - Trigger execution can be disabled before such actions.
- Other risks with triggers:
 - Error leading to failure of critical transactions that set off the trigger
 - Cascading execution



Recursive Queries



Recursion in SQL

- SQL:1999 permits recursive view definition

Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
-- prepare our existing table
-- INSERT INTO course VALUES ('CS-401', 'Data mining', 'Comp. Sci.', 3);
-- INSERT INTO prereq VALUES ('CS-401', 'CS-347');
```

```
SELECT * FROM prereq
CONNECT BY PRIOR prereq_id = course_id
START WITH course_id = 'CS-401';
      OR

WITH    required (course_id, prereq_id) AS
      (
        SELECT  course_id, prereq_id
        FROM    prereq
        WHERE    course_id = 'CS-401'
        UNION ALL
        SELECT  required.course_id, prereq.prereq_id
        FROM    required
        JOIN     prereq
        ON       required.prereq_id = prereq.course_id
      )
SELECT  *
FROM    required;
```



The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
 - ▶ This can give only a fixed number of levels of managers
 - ▶ Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work
 - ▶ Alternative: write a procedure to iterate as many times as required
 - See procedure *findAllPrereqs* in book



Advanced Aggregation Features



Advanced Aggregation

```
-- prepare table
-- create table
CREATE TABLE student_grades(
  student_id numeric(8, 0),
  student_name nvarchar2(64),
  dept_name nvarchar2(32),
  cgpa numeric(4, 2)
);

-- insert values
INSERT ALL
INTO student_grades VALUES (1, 'Tom', 'CSE', 3.4)
INTO student_grades VALUES (2, 'Leo', 'CSE', 3.3)
INTO student_grades VALUES (3, 'Chris', 'CSE', 3.6)
INTO student_grades VALUES (4, 'Michael', 'EEE', 3.8)
INTO student_grades VALUES (5, 'Quentin', 'EEE', 3.6)
INTO student_grades VALUES (6, 'Matt', 'EEE', 3.5)
INTO student_grades VALUES (7, 'Jerry', 'CE', 3.1)
INTO student_grades VALUES (8, 'David', 'CE', 3.3)
INTO student_grades VALUES (9, 'Jason', 'CE', 2.8)
INTO student_grades VALUES (10, 'Dwayne', 'CE', 3.2)
INTO student_grades VALUES (11, 'Richard', 'CSE', 3.4)
SELECT 1 FROM dual;
```



Ranking

- Ranking is done in conjunction with an order by specification.
- Suppose we are given a relation
student_grades(ID, GPA)
giving the grade-point average of each student
- Find the rank of each student.

```
SELECT student_id,  
       student_name,  
       cgpa,  
       dept_name,  
       rank() OVER (ORDER BY cgpa desc) AS student_rank,  
       dense_rank() OVER (ORDER BY cgpa desc) AS student_dense_rank  
FROM student_grades;
```

- Ranking may leave gaps: e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
 - **dense_rank** does not leave gaps, so next dense rank would be 2



Ranking (Cont.)

- Ranking can be done within partition of the data.
- “Find the rank of students within each department.”

```
SELECT student_id,  
       student_name,  
       cgpa,  
       dept_name,  
       rank() OVER (PARTITION BY dept_name ORDER BY cgpa desc) AS student_rank,  
       dense_rank() OVER (PARTITION BY dept_name ORDER BY cgpa desc) AS  
       student_dense_rank  
FROM student_grades;
```

- Multiple **rank** clauses can occur in a single **select** clause.
- Ranking is done *after* applying **group by** clause/aggregation
- Can be used to find top-n results
 - More general than the **limit** *n* clause supported by many databases, since it allows top-n within each partition



Exercise

- Find the student who is ranked second in his/her department.



Exercise

- Find the student who is ranked second in his/her department.

```
WITH student_ranks AS (  
    SELECT student_id,  
           student_name,  
           cgpa,  
           dept_name,  
           rank() OVER (PARTITION BY dept_name ORDER BY cgpa desc) AS student_rank,  
           dense_rank() OVER (PARTITION BY dept_name ORDER BY cgpa desc) AS student_dense_rank  
    FROM student_grades  
)  
SELECT *  
FROM student_ranks  
WHERE student_rank = 2;
```



Ranking (Cont.)

- Other ranking functions:
 - **percent_rank** (within partition, if partitioning is done)
 - **cume_dist** (cumulative distribution)
 - fraction of tuples with preceding values
 - **row_number** (non-deterministic in presence of duplicates)

SQL:1999 permits the user to specify **nulls first** or **nulls last**

```
select ID,  
       rank ( ) over (order by cgpa desc nulls last) as s_rank  
from student_grades
```



Ranking (Cont.)

- For a given constant n , the ranking the function $ntile(n)$ takes the tuples in each partition in the specified order, and divides them into n buckets with equal numbers of tuples.
- E.g.,

```
select student_id,  
       student_name,  
       dept_name,  
       cgpa,  
       ntile(4) over (order by cgpa desc) as quartile  
from student_grades;
```



Windowing

- Used to smooth out random variations.
- E.g., **moving average**: “Given sales values for each date, calculate for each date the average of the sales on that day, the previous day, and the next day”
- **Window specification** in SQL:
 - Given relation emp(*empno*, *ename*, *job*, *mgr*, *hiredate*, *sal*, *comm*, *deptno*)

```
SELECT ename,  
       sal,  
       sum(sal) OVER (PARTITION BY deptno order by sal) AS running_salary  
FROM emp;
```

```
SELECT empno,  
       ename,  
       job,  
       sal,  
       sum(sal) OVER (PARTITION BY deptno ORDER BY sal DESC ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING),  
       deptno  
FROM emp;
```

```
SELECT empno,  
       ename,  
       job,  
       sal,  
       sum(sal) OVER (PARTITION BY deptno ORDER BY sal DESC RANGE BETWEEN 1 PRECEDING AND 1 FOLLOWING),  
       deptno  
FROM emp;
```



Windowing

- Examples of other window specifications:
 - **between rows unbounded preceding and current**
 - **rows unbounded preceding**
 - **range between 10 preceding and current row**
 - All rows with values between current row value -10 to current value
 - **range interval 10 day preceding**
 - Not including current row



Windowing (Cont.)

- Can do windowing within partitions
- E.g., Given a relation *transaction* (*account_number*, *date_time*, *value*), where *value* is positive for a deposit and negative for a withdrawal
 - “Find total balance of each account after each transaction on the account”

```
select account_number, date_time,  
       sum (value) over  
         (partition by account_number  
          order by date_time  
          rows unbounded preceding)  
       as balance  
from transaction  
order by account_number, date_time
```



Windowing (Cont.)

- Lead/Lag

```
SELECT empno,  
       ename,  
       job,  
       sal,  
       LAG(sal, 1, 0) OVER (PARTITION BY deptno ORDER BY sal desc) AS sal_prev,  
       sal - LAG(sal, 1, 0) OVER (PARTITION BY deptno ORDER BY sal desc) AS sal_diff,  
       deptno  
FROM   emp;
```



End of Chapter 5