

# Structure of Processes

CSI 309: Operating System Concepts  
United International University

# PROCESSES AND PROGRAMS

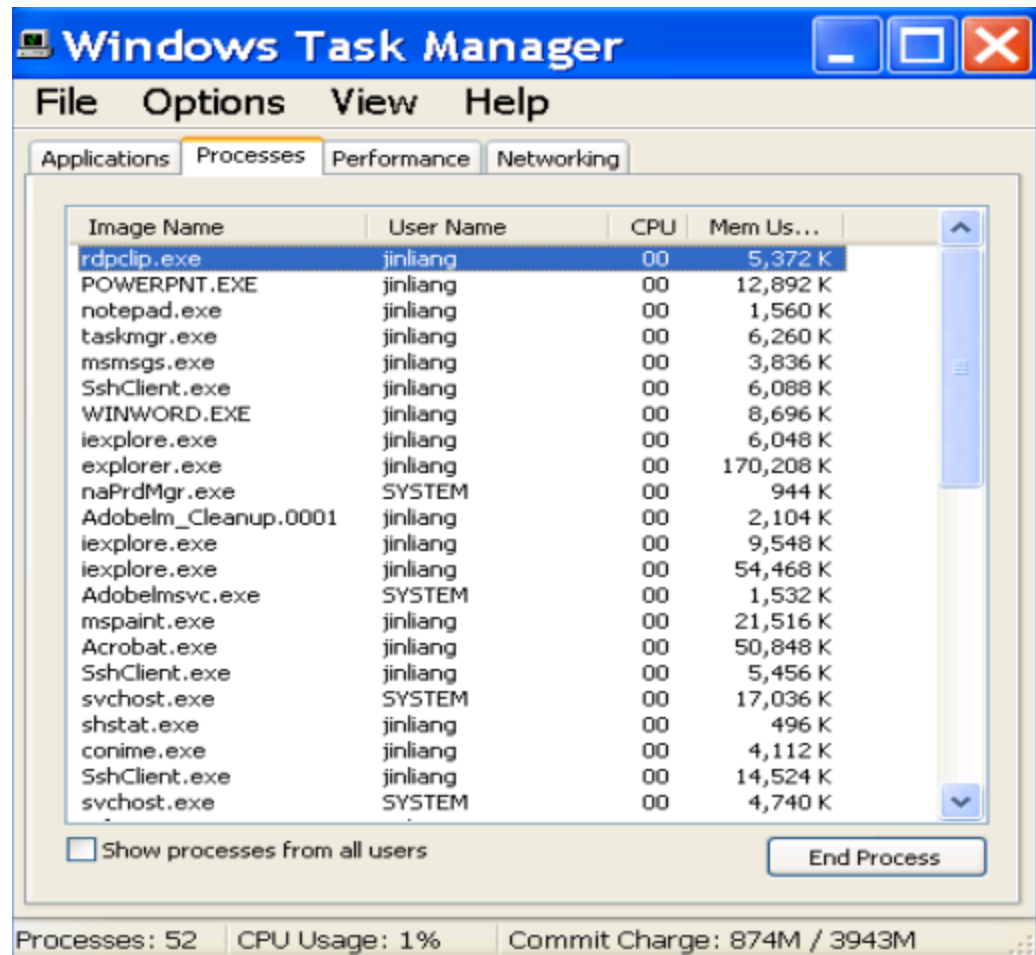
# Users, Programs, Processes

- Users have accounts on the system
- Users launch programs
  - Many users may launch same program
  - One user may launch many instances of the same program
- Program: an algorithm expressed in some suitable notation.
- Process: a program in execution.

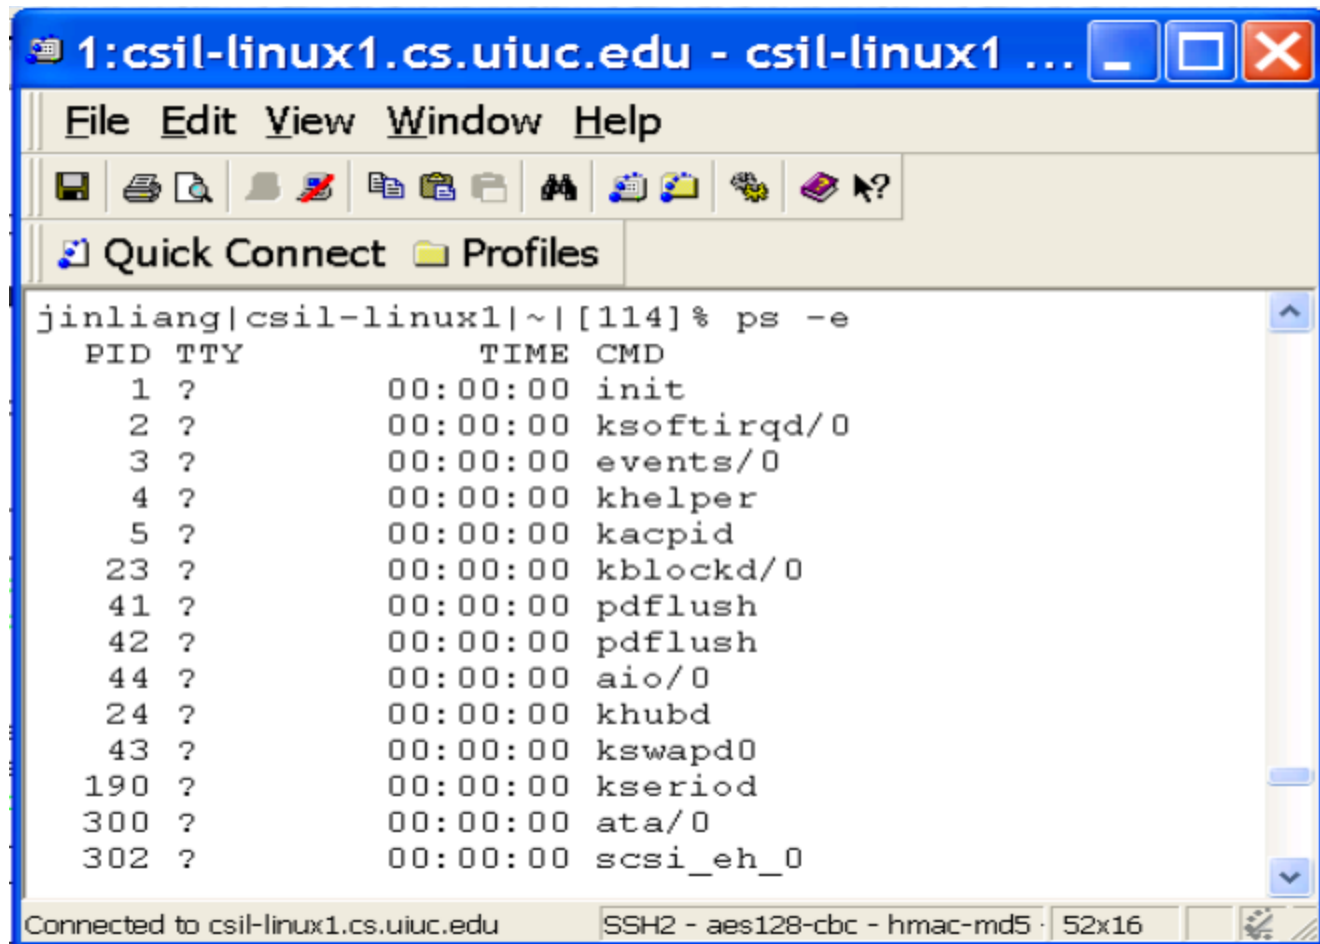
# What is a process?

- A task created by the OS, running in a restricted virtual machine environment –a virtual CPU, virtual memory environment, interface to the OS via system calls
- The unit of execution
- Operating system provided ***abstraction*** to represent what is **needed** to run a single **program**
- The same as “job” or “task” or “sequential process”.

# Example: Windows Task Manager



# Example: ps in Unix



The screenshot shows a terminal window titled "1:csil-linux1.cs.uiuc.edu - csil-linux1 ...". The window has a menu bar with "File", "Edit", "View", "Window", and "Help". Below the menu bar is a toolbar with various icons. The main area of the window displays the output of the command "ps -e" executed by the user "jinliang". The output is a table with four columns: "PID", "TTY", "TIME", and "CMD". The table lists system processes including "init", "ksoftirqd/0", "events/0", "khelper", "kacpid", "kblockd/0", "pdflush", "aio/0", "khubd", "kswapd0", "kseriod", "ata/0", and "scsi\_eh\_0". The status bar at the bottom indicates the connection is to "csil-linux1.cs.uiuc.edu" and shows encryption details: "SSH2 - aes128-cbc - hmac-md5" and "52x16".

```
jinliang|csil-linux1|~|[114]% ps -e
```

PID	TTY	TIME	CMD
1	?	00:00:00	init
2	?	00:00:00	ksoftirqd/0
3	?	00:00:00	events/0
4	?	00:00:00	khelper
5	?	00:00:00	kacpid
23	?	00:00:00	kblockd/0
41	?	00:00:00	pdflush
42	?	00:00:00	pdflush
44	?	00:00:00	aio/0
24	?	00:00:00	khubd
43	?	00:00:00	kswapd0
190	?	00:00:00	kseriod
300	?	00:00:00	ata/0
302	?	00:00:00	scsi_eh_0

Connected to csil-linux1.cs.uiuc.edu SSH2 - aes128-cbc - hmac-md5 52x16

# What is a program?

- A Program is an **executable file** that contains:
  - **Code:** Machine instructions
  - **Data:** Variables stored and manipulated in memory
    - initialized variables (globals)
    - dynamically allocated variables (malloc, new)
    - stack variables (C automatic variables, function arguments)
  - **DLL:** libraries that were not compiled or linked with the program
    - containing code & data, possibly shared with other programs

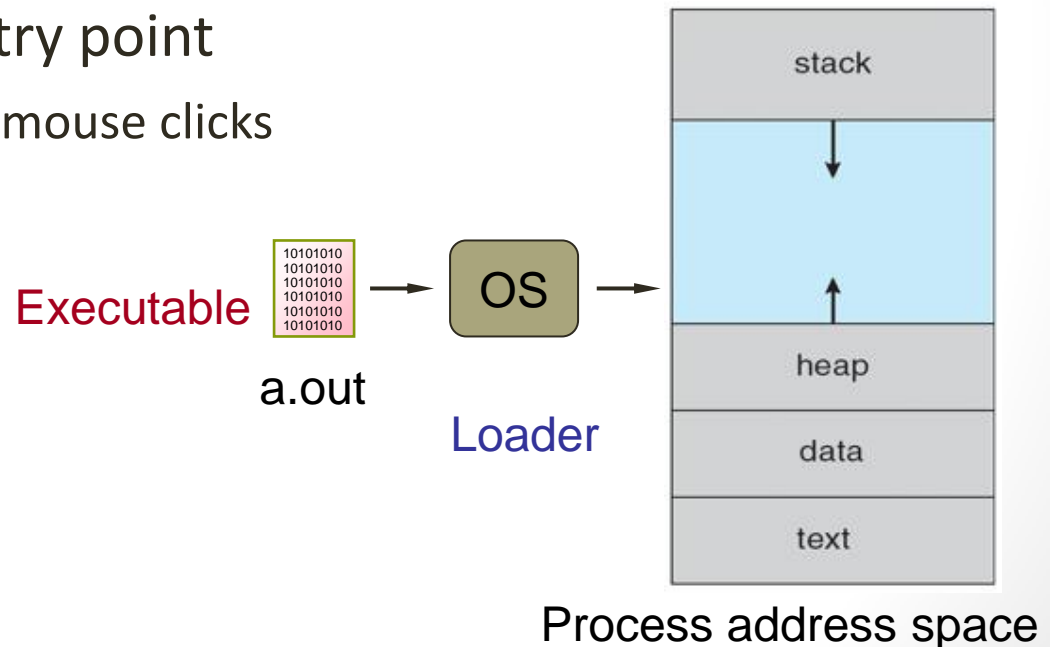
# Process != Program

- A process is an executing program.
- Example:
  - We can run 2 instances of Mozilla Firefox:
    - Same program
    - Separate processes
- **Program** is passive: Code + data
- **Process** is running program: stack, registers, program counter




# How Program Becomes a Process

- When a program is launched
  - OS **loads** program into memory
  - Creates kernel data structure for the process
  - **Initializes** data
  - Starts from an entry point
    - e.g., main(), GUI mouse clicks



# Process address space


- set of all memory addresses accessible by a process



# MULTIPROGRAMMING


# Multiprogramming

- Each process has its own address space (virtual memory address)
- Even if two processes are running the same program, they have their own address space
- OS schedules processes




# Multiprogramming

One program counter




(a)

Four program counters



(b)



(c)

- Multiprogramming of **4** processes in a **single** CPU  
CPU switches from one process to other process
- Only 1 **physical** program counter  
4 **logical** program counters
- Conceptual model of 4 independent, sequential processes
- Only **1** program active at any instant.
- Real life analogy?

**A daycare teacher trying to feed 4 infants.**

# PROCESS OPERATIONS

Process Creation

Process Termination

Process State Transitions

# Process Creation

- System initialization
  - Boot, reboot.
- Execution of a process creation system call
  - `fork()`
- User request to create a new process
  - Command line or click an icon.
- Initiation of a batch job

# Process Termination

- Normal exit (voluntary)
  - End of main()
- Error exit (voluntary)
  - exit(2)
- Fatal error (involuntary)
  - Divide by 0, core dump
- Killed by another process (involuntary)
  - kill proclD, end task

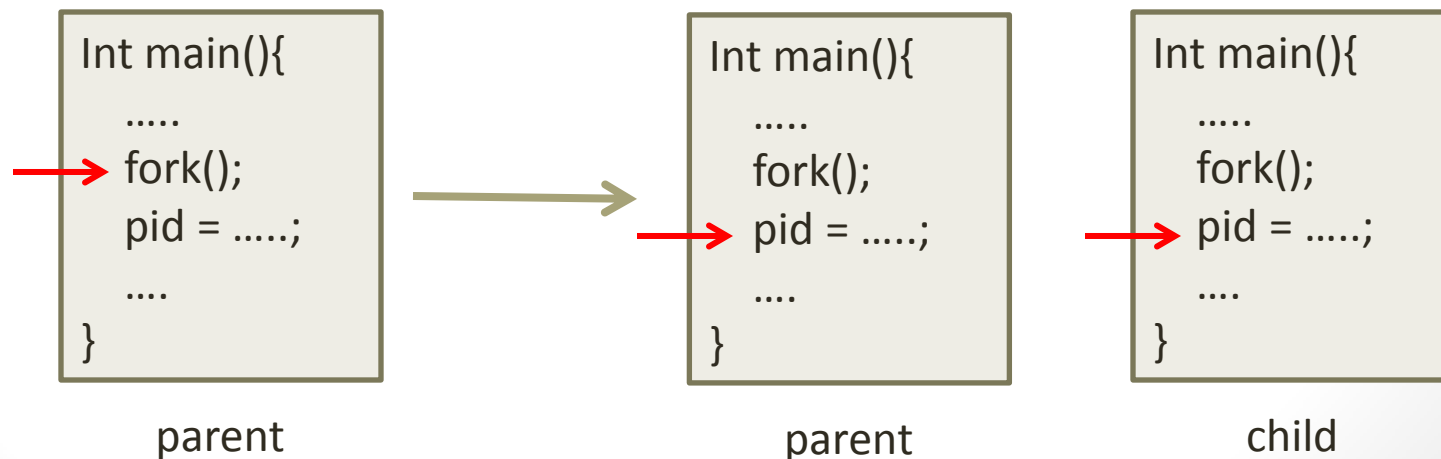


# Example: fork() in Unix

- `int childpid;`
- ```
if ((childpid = fork()) == -1) {  
    perror("can't create a new process");  
    exit(1);  
} else if (childpid == 0) {  
    // executes child process code  
    exit(0);  
} else {  
    // executes parent process code  
    exit(0);  
}
```

# The fork() System Call

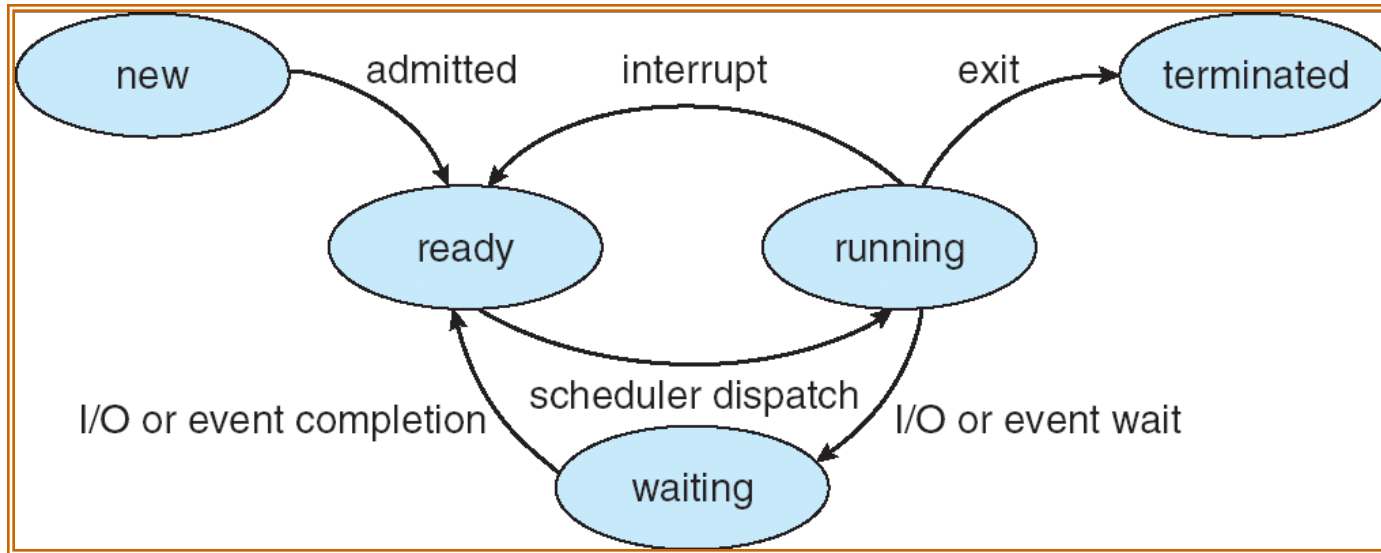
- When fork() is called in process A
  - Control is switched to kernel
  - Kernel creates new process B by copying A's
    - Address space
    - Kernel data structure (process descriptor)
  - OS now has two identical processes to run. Both resume from after fork(). However, return value of fork() will be different.



# Process States

- Many processes in system, only one on CPU
- “Execution State” of a process:
  - Indicates what it is doing
  - Basically 3 states:
    1. Ready: waiting to be assigned to the CPU
    2. Running: executing instructions on the CPU
    3. Waiting: waiting for an event, e.g. I/O completion
- Process moves across different states

# Process State Transitions



- As a process executes, it changes *state*
  1. **new**: The process is being created
  2. **ready**: The process is waiting to run
  3. **running**: Instructions are being executed
  4. **waiting**: Process waiting for some event to occur
  5. **terminated**: The process has finished execution

# Process State Transitions

Processes hop across states as a result of:

- Actions they perform, e.g. system calls
- Actions performed by OS, e.g. rescheduling
- External actions, e.g. I/O


# Process State Transitions

- **Running → Block**: process discovers that it cannot continue. If running process initiates an I/O operation before its allotted time expires, the running process voluntarily relinquishes the CPU.
- **Running → Ready**: scheduler decides that the running process has run long enough and it is time to let another process have CPU time.
- **Ready → Running**: all other processes have had their share and it is time for the first process to run again
- **Blocked → Ready**: external event for which a process was waiting (such as arrival of input) happens.
- **New → Ready** : process is created.
- **Running → Terminated** : process has finished execution.

# PROCESS HIERARCHIES AND PROCESS DATA STRUCTURES

# Process Hierarchies

- When process A creates process B, A is called the “parent” process, B is the “child”
- Forms a hierarchy
  - UNIX calls this a “process group”
  - A special process present in boot image is **init**
- **Windows** has **no** process hierarchy
  - Processes are independent after creation





# Process Data Structures

- OS represents a process using a PCB
  - Process Control Block
  - Has all the details of a process

|                                  |                             |
|----------------------------------|-----------------------------|
| <b>Process Id</b>                | <b>Security Credentials</b> |
| <b>Process State</b>             | <b>Username of owner</b>    |
| <b>General Purpose Registers</b> | <b>Queue Pointers</b>       |
| <b>Stack Pointer</b>             | <b>Signal Masks</b>         |
| <b>Program Counter</b>           | <b>Memory Management</b>    |
| <b>Accounting Info</b>           | <b>...</b>                  |

# Process Control Block (PCB)

## Fields of a Process Table Entry

| <b>Process management</b>                                                                                                                                                                                                                                                         | <b>Memory management</b>                                                       | <b>File management</b>                                                         |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|--------------------------------------------------------------------------------|
| Registers<br>Program counter<br>Program status word<br>Stack pointer<br>Process state<br>Priority<br>Scheduling parameters<br>Process ID<br>Parent process<br>Process group<br>Signals<br>Time when process started<br>CPU time used<br>Children's CPU time<br>Time of next alarm | Pointer to text segment<br>Pointer to data segment<br>Pointer to stack segment | Root directory<br>Working directory<br>File descriptors<br>User ID<br>Group ID |

# Reference

Modern Operating Systems

Section 2.1